

RAG-Guardrails Integration for AI Content Control

Rakesh More*

A J Gallagher USA (Independent Researcher)

Chicago, IL, USA

rakeshmore@gmail.com

Abstract

Generative AI, particularly large language models (LLMs), has shown remarkable potential across domains such as healthcare, legal services, and finance. However, their adoption is hindered by two persistent challenges: hallucination-where models generate factually incorrect information-and the risk of producing biased or unsafe content. This paper proposes a hybrid framework that integrates Retrieval-Augmented Generation (RAG) with NVIDIA NeMo Guardrails to address these concerns. RAG mitigates hallucinations by grounding model outputs in externally retrieved, trusted data sources, while NeMo Guardrails enforce domain-specific safety and compliance constraints through predefined behavioral policies. Empirical evaluations demonstrate that this combined approach reduces hallucinated content by 30–45% and improves safety and policy adherence across multiple enterprise use cases. The system exhibits strong potential for deployment in regulated, high-stakes environments. Future work will focus on enhancing real-time responsiveness and expanding multilingual and culturally adaptive capabilities. The proposed framework offers a scalable foundation for building trustworthy, domain-aligned generative AI solutions.

CCS Concepts

• **Security and privacy** → Software and application security; Domain-specific security and privacy architectures.

Keywords

Large language models, generative AI, hallucination, RAG, NeMo Guardrails, content safety, responsible AI, retrieval-augmented generation

ACM Reference Format:

Rakesh More. 2025. RAG-Guardrails Integration for AI Content Control. In *The 18th International Conference on Computer Science and Information Technology (ICCSIT 2025)*, October 27–29, 2025, Paris, France. ACM, New York, NY, USA, 9 pages. <https://doi.org/10.1145/3783862.3783896>

1 Introduction

Large language models-like GPT-4 from OpenAI [4], Google’s PaLM, and Meta’s LLaMA-have really shaken up how we approach language tasks on computers. Built on top of enormous collections of text, these models use complicated architecture (transformers, mostly) to make eerily fluent language predictions. As a result,

we’re relying on them more than ever, whether it’s powering virtual assistants, helping automate business summaries [1], delivering tutoring experiences, or answering random questions with speed and confidence. Benchmark studies consistently show that they excel at generating text [6], handling translation [7], and summarizing documents [1]. Despite all the hype and impressive stats, a persistent issue remains hallucinations. It’s not a science fiction problem-just a term for when the model confidently “makes stuff up” that sounds convincing but isn’t grounded or sometimes can’t be verified at all [15]. This quickly goes from quirky to potentially dangerous if you imagine an LLM writing up a medical answer that’s wrong [8] or producing a news summary that misleads readers [9]. Even in research, users have caught models fabricating references that look real but aren’t tied to any actual publication [15]. Another minefield: unsafe or non-compliant outputs. Without guardrails, these language models can blurt out responses that are biased, offensive, or cross legal boundaries. This poses a major challenge for deploying these tools in fields like healthcare, law, or education, where accuracy and sensitivity are essential.

To address these risks, I propose a straightforward yet promising hybrid strategy. This approach integrates RAG with NVIDIA NeMo Guardrails to create a comprehensive solution addressing both hallucination and safety concerns. Our approach leverages the complementary strengths of both methodologies: RAG provides external knowledge grounding to reduce hallucinations, while NeMo Guardrails ensures comprehensive safety enforcement across all interaction phases.

The key contributions of this work include:

- **Integration Architecture:** A systematic framework for combining RAG and NeMo Guardrails that addresses the limitations of individual approaches while maximizing their complementary benefits.
- **Comprehensive Evaluation:** Systematic assessment across multiple domains (healthcare, legal, finance) using established benchmarks and novel evaluation metrics specifically designed for hybrid systems.
- **Security and Privacy Analysis:** Detailed examination of security implications, privacy protection mechanisms, and compliance considerations for enterprise deployment.
- **Deployment Guidelines:** Practical recommendations for implementing the integrated framework in production environments, including performance optimization and scalability considerations.

Our experimental results demonstrate significant improvements in both safety and accuracy metrics. The integrated system achieves 97% accuracy in detecting safety violations while reducing hallucination rates by 32-47% across various domains. Importantly, these improvements are achieved with minimal latency overhead (< 500ms), making the solution viable for real-time applications.

*Corresponding author.



This work is licensed under a Creative Commons Attribution 4.0 International License. *ICCSIT 2025, Paris, France*

© 2025 Copyright held by the owner/author(s).

ACM ISBN 979-8-4007-1858-8/2025/10

<https://doi.org/10.1145/3783862.3783896>

2 BACKGROUND

2.1 AI Hallucinations

AI hallucinations represent a critical phenomenon where generative models produce factually incorrect information while maintaining apparent confidence in their outputs. These manifestations range from minor inaccuracies to completely fabricated content, including non-existent citations and sources. The term "hallucination" aptly describes AI's tendency to perceive patterns or relationships that do not exist in reality, resulting in plausible sounding but factually incorrect outputs [5]. The primary causes of hallucinations stem from several interconnected factors. Training data quality issues, including incomplete or biased datasets, contribute significantly to this problem. Model overfitting, where systems memorize training examples rather than learning generalizable patterns, represents another major contributor. Additionally, current AI architecture lacks true factual understanding [10], relying instead on statistical pattern recognition, which can lead to misinterpretation of complex information relationships.

The consequences of AI hallucinations are particularly severe in critical applications. In healthcare settings, incorrect diagnostic suggestions may lead to inappropriate treatment decisions [14]. Financial applications face risks of erroneous predictions resulting in substantial economic losses. Furthermore, hallucinations contribute to misinformation propagation, undermining public trust in AI systems. The inconsistent nature of generative output also poses challenges for scientific reproducibility, a fundamental requirement in research applications.

2.2 Security and Privacy Risks

Generative AI systems face multifaceted security challenges that extend beyond traditional cybersecurity concerns. The primary threat vector involves prompt injection attacks, where malicious actors manipulate AI systems through carefully crafted inputs designed to bypass security measures or extract sensitive information. Prompt injection manifests in two primary forms: direct injection through immediate user prompts containing hidden instructions, and indirect injection involving long-term system compromise through poisoned data sources. The latter represents a particularly insidious threat as it corrupts the AI's knowledge base over time, leading to persistent vulnerabilities that are difficult to detect and remediate.

Privacy concerns constitute another significant risk category. Many generative AI systems collect and store user interactions for model improvement purposes, creating potential exposure points for personally identifiable information (PII). This data collection practice raises compliance concerns with privacy regulations and creates risks for unintended information disclosure. System-level vulnerabilities further exacerbate these challenges. Poorly secured APIs, plugin architectures, and hosting environments can serve as attack vectors. Furthermore, adversarial attacks involving subtle input modifications designed to manipulate outputs represent sophisticated threats that can circumvent traditional security measures.

2.3 Traditional Mitigation Approaches

Addressing these challenges requires comprehensive, multi-layered approaches. For hallucination prevention, effective strategies include domain restriction to well-defined areas, rigorous training data curation, template-based generation to constrain outputs, and integrated feedback mechanisms for continuous improvement. Security risk mitigation requires defense-in-depth strategies that encompass input filtering through multiple validation stages, output verification before deployment, access control based on the principle of least privilege, and mandatory human oversight for sensitive operations. Operational security measures include system isolation, continuous monitoring, regular adversarial testing, and comprehensive user education programs.

Table 1 summarizes common generative AI challenges such as hallucination and security risks, along with corresponding mitigation strategies derived from prior studies.

3 RELATED WORK

3.1 Retrieval-Augmented Generation (RAG) - Enhancing through External Knowledge Integration

Retrieval-Augmented Generation represents a paradigm shift in large language model architecture, addressing the inherent limitations of models that rely exclusively on parametric knowledge acquired during training. Traditional language models operate within the constraints of their training data, potentially generating responses based on outdated or incomplete information. RAG systems overcome this limitation by incorporating dynamic access to external knowledge sources during the generation process [2].

The fundamental architecture of RAG systems consists of two primary components: a retrieval mechanism and a generation model. When processing a query, the system first converts the input into a dense vector representation through embedding techniques. This embedding facilitates semantic search across a vectorized knowledge base, identifying the most relevant documents or passages. The retrieved information is then concatenated with the original query and fed into the language model, which generates responses grounded in both parametric knowledge and external evidence [16]. RAG pairs a retriever with a generator: a query is embedded and used to retrieve documents from a vector store; retrieved documents are provided to the generator for grounded output [2]. Benchmarking studies [12] have systematically evaluated the factuality and retrieval efficiency of RAG architectures, revealing the importance of contextual grounding. Recent improvements, such as Iterative Retrieval-Generation Synergy [13] refine retrieval loops to enhance factual consistency and reduce error propagation. Domain-specific applications, such as CarExpert [11], further illustrate RAG's adaptability across specialized environments, including automotive and enterprise contexts.

Advanced RAG implementations employ iterative retrieval strategies, where the system performs multiple retrieval cycles to gather comprehensive information. This approach proves particularly effective for complex queries requiring synthesis of information from multiple sources. The iterative process continues until sufficient context is obtained or a predetermined threshold is reached.

Table 1: Generative AI challenges and general mitigation strategies

Challenge Category	Specific Manifestations/ Examples	General Mitigation Strategies
AI Hallucinations	Incorrect predictions, False positives/negatives, fabricated links/citations	High-quality/relevant training data, limiting outcomes/responses, Data templates, providing explicit feedback
Sensitive Content/Security Risks	Prompt injection, PII leakage, Biased/skewed outputs, Infrastructure vulnerabilities, Inappropriate content	Constraining model behavior, Input/Output filtering, least privilege, Human oversight, Red teaming, Continuous monitoring, and User education


**Figure 1: Retrieval-Augmented Generation (RAG) Architecture**

Figure 1 shows the Retrieval-Augmented Generation (RAG) architecture, illustrating how external knowledge is retrieved and integrated into the generation process. The diagram highlights the retriever and generator components and the flow of information between them.

Addressing Hallucination Through External Grounding

The generation of plausible yet factually incorrect information represents a critical challenge for deployment in high-stakes applications. RAG systems mitigate this issue by anchoring responses in verifiable external sources, providing explicit provenance for generated content. This grounding mechanism enables users to trace the origin of information and verify claims against source documents. However, RAG systems are not immune to accuracy challenges. Conflicts may arise between the model’s parametric knowledge and retrieved information, potentially leading to inconsistent or contradictory responses. Additionally, the system may misinterpret retrieved content or exhibit selection bias toward information that aligns with pre-existing model biases. Recent developments in RAG architecture, including Decision-based Self-Supervised Pre-training for RAG (DSSP-RAG) and Iterative Retrieval-Generation (ITER-RETGEN), address these limitations through enhanced decision-making mechanisms. These approaches incorporate quality assessment modules that evaluate when external retrieval is necessary and implement filtering mechanisms to exclude unreliable sources.

Security and Privacy Considerations - The integration of external knowledge sources introduces novel security vulnerabilities that extend beyond traditional model-based risks. RAG systems inherit the security posture of their underlying knowledge bases, making them susceptible to data poisoning attacks where malicious actors introduce false information into external sources. This vulnerability is particularly concerning given the authoritative nature of externally sourced information in user perception. Privacy concerns arise from the dual nature of RAG systems, which process both user queries and potentially sensitive external documents.

The retrieval process may inadvertently expose confidential information, while the generation component may leak sensitive data from the knowledge base. These risks necessitate comprehensive privacy-preserving mechanisms throughout the RAG pipeline. Effective RAG security requires a multi-layered approach encompassing source verification, content filtering, and output monitoring. Organizations must implement robust authentication mechanisms for knowledge sources, continuous monitoring for content integrity, and regular security assessments to identify potential vulnerabilities. The development of secure RAG systems demands consideration of attack vectors from system design through deployment.

3.2 NVIDIA NeMo Guardrails - Comprehensive AI Safety Framework

NVIDIA NeMo Guardrails provides a systematic approach to AI safety through programmable constraints that monitor and control language model behavior across multiple interaction phases. Unlike post-hoc safety measures that focus solely on output filtering, NeMo Guardrails implements a comprehensive safety architecture that governs the entire conversation flow [3].

Multi-Rail Safety Architecture - The NeMo Guardrails framework employs a multi-layered safety approach through five distinct rail categories, each addressing specific aspects of AI interaction safety:

- **Input Rails:** These components intercept and analyze user inputs before processing, identifying potentially harmful content, personally identifiable information (PII), or attempts at prompt injection. Input rails serve as the first line of defense against malicious or inappropriate queries.
- **Dialog Rails:** These mechanisms maintain conversational coherence and ensure adherence to predefined conversation policies. Dialog rails prevent topic drift, enforce business rules, and maintain appropriate interaction boundaries.




Figure 2: NeMo Guardrails Multi-Rail Architecture

- **Retrieval Rails:** Specifically designed for RAG-enabled systems, these components filter and validate external content before integration into the generation process. Retrieval rails assess source credibility, content relevance, and potential security risks.
- **Execution Rails:** These controls govern the system’s ability to execute external actions, preventing unauthorized code execution or access to restricted resources. Execution rails are particularly critical for AI systems with tool-calling capabilities.
- **Output Rails:** The final safety checkpoint, these components analyze generated responses for hallucinations, harmful content, policy violations, or sensitive information leakage before delivery to users.

Figure 2 shows the NeMo Guardrails multi-rail architecture, emphasizing how input, dialog, retrieval, execution, and output rails collaborate to enforce safety and policy constraints throughout the model’s lifecycle.

Advanced Hallucination Detection - NeMo Guardrails incorporates sophisticated hallucination detection mechanisms that extend beyond simple content filtering. The system employs self-consistency checking, where the model is prompted to validate its own responses against established facts. Additionally, external validation tools such as AlignScore provide independent fact-checking capabilities. For RAG-specific applications, NeMo Guardrails integrates Patronus Lynx, a specialized tool designed to detect hallucinations that arise from the misintegration of retrieved information. This capability addresses a critical gap in RAG systems, where hallucinations may result from the improper synthesis of multiple information sources rather than purely generative errors. NeMo Guardrails defines a multi-rail safety architecture (input, dialog, retrieval, execution, output) which enforces policies through programmable rails, integrates with tools like Presidio for PII detection, and supports hallucination detection (e.g., Patronus Lynx) and auditing. Healthcare-specific extensions [14] demonstrate how domain-oriented guardrails enhance compliance and policy adherence in regulated sectors.

Privacy Protection and Compliance - The framework implements comprehensive privacy protection through integration with Microsoft Presidio, which provides automated detection and anonymization of sensitive information across multiple data types. This protection extends beyond output filtering to include input

sanitization and retrieval content processing. NeMo Guardrails supports regulatory compliance through configurable policies that can be tailored to specific industry requirements. The system provides audit trails and logging capabilities essential for demonstrating compliance with data protection regulations such as GDPR and HIPAA.

Integration and Extensibility - Rather than replacing existing safety tools, NeMo Guardrails provides a unified orchestration layer that integrates with established safety solutions, including LlamaGuard, ActiveFence, and OpenAI’s moderation API. This approach enables organizations to leverage existing investments while benefiting from centralized safety management. The framework’s extensibility through the Colang domain-specific language allows organizations to implement custom safety policies and conversation flows tailored to specific use cases. This flexibility ensures that safety measures can adapt to evolving requirements without requiring system redesign.

Performance and Scalability - Despite its comprehensive safety mechanisms, NeMo Guardrails maintains operational efficiency with minimal latency overhead (typically under 500 milliseconds). The system’s architecture enables selective rail activation based on risk assessment, allowing organizations to balance safety requirements with performance considerations.

The framework represents a significant advancement in AI safety methodology, transitioning from reactive safety measures to proactive, integrated safety architectures. This approach demonstrates that comprehensive AI safety can be achieved without compromising system performance or user experience, facilitating responsible AI deployment at enterprise scale.

4 Methodology

Despite the advances offered by individual methodologies, neither guardrail systems nor retrieval augmentation methods alone are sufficient to fully mitigate hallucination risks and security vulnerabilities. Guardrails can enforce rules and reduce harmful output, but they may not always capture nuanced factual inaccuracies. Conversely, retrieval systems can enhance fact-checking but may lack stringent enforcement of ethical policies. Our hybrid framework is motivated by the need to conjoin these complementary approaches, ensuring that safety is embedded in both the generation process and its factual underpinning.

Table 2 presents cross-domain evaluation results showing that the integrated RAG-Guardrails framework improves accuracy, safety, and privacy across healthcare, legal, and financial domains..

4.1 Formalized RAG-Guardrails Integration Architecture


Our hybrid framework establishes a formal interoperability specification between retrieval validation and guardrail enforcement phases through a deterministic state-machine-driven coordination schema. The integration operates through five formally defined phases with explicit transition semantics:

- State-Machine Coordination Protocol

The integration follows a finite state automaton $S = (Q, \Sigma, \delta, q_0, F)$ where:

Table 2: Comparative analysis of existing approaches

Approach	Hallucination Mitigation	Safety Enforcement	Privacy Protection	Deployment Complexity	Performance Overhead
Baseline LLM	Low	Low	Low	Low	Low
RAG Only	High	Low	Medium	Medium	Medium
Guardrails Only	Medium	High	High	Medium	Low
Fine-tuning	Medium	Medium	Low	High	Low
Prompt Engineering	Medium	Medium	Low	Low	Low
Our Approach	High	High	High	Medium	Medium

**Figure 3: Integrating RAG and NeMo Guardrails System architecture diagrams**

- $Q = \{\text{INPUT, RETRIEVAL, SYNTHESIS, GENERATION, OUTPUT}\}$ represents system states
- $\Sigma = \{\text{user_query, safety_check, document_retrieval, response_generation}\}$ represents input alphabet
- $\delta: Q \times \Sigma \rightarrow Q$ defines deterministic state transitions
- $q_0 = \text{INPUT}$ represents the initial state
- $F = \{\text{OUTPUT, SAFETY_VIOLATION}\}$ represents final states

Transition Semantics:

- $\delta(\text{INPUT, user_query}) = \text{RETRIEVAL}$ if $\text{safety_check}(\text{input}) = \text{SAFE}$
- $\delta(\text{INPUT, user_query}) = \text{SAFETY_VIOLATION}$ if $\text{safety_check}(\text{input}) = \text{UNSAFE}$
- $\delta(\text{RETRIEVAL, document_retrieval}) = \text{SYNTHESIS}$ if $\text{validation_check}(\text{docs}) = \text{VALID}$
- $\delta(\text{SYNTHESIS, response_generation}) = \text{GENERATION}$ if $\text{context_integration} = \text{SUCCESS}$
- $\delta(\text{GENERATION, safety_check}) = \text{OUTPUT}$ if $\text{all_rails_passed} = \text{TRUE}$

- Cross-Layer Protocol Specification

Data Provenance Handling Protocol:

- **Input Layer (I):** $P_I = \{\text{query_id, user_context, safety_meta-data}\}$

- **Retrieval Layer (R):** $P_R = \{\text{document_id, source_credibility, retrieval_confidence}\}$
- **Output Layer (O):** $P_O = \{\text{response_id, factual_grounding, safety_compliance}\}$

Verifiable Consistency Framework: Each data transformation maintains a provenance chain: $P_I \rightarrow P_R \rightarrow P_O$ with cryptographic integrity verification.

Figure 3 shows the integrated RAG–Guardrails architecture combines retrieval-based factual grounding with programmable guardrails that monitor and regulate interactions at each system stage.

The diagram illustrates a five-phase data flow pipeline: Input, Retrieval, Synthesis, Generation, and Output. Throughout the process, multiple safety rails—input, retrieval, and output—continuously monitor for hallucinations, privacy risks, and policy violations. Arrows between modules indicate bidirectional information exchange and provenance tracking, while a control layer enforces deterministic state transitions and conflict resolution. The architecture emphasizes how combining RAG with NeMo Guardrails ensures factual accuracy, data privacy, and reliable AI behavior.

4.2 Multi-Rail Safety Framework Implementation

- Rail-Specific Operational Semantics

- Input Rails (R_input):

```
def input_rail_processing(query):
    return {
        'jailbreak_detection': heuristic_analysis(query),
        'topic_validation': topical_matching(query),
        'pii_detection': presidio_analysis(query),
        'safety_score': composite_safety_evaluation(query)
    }
```

- Retrieval Rails (R_retrieval):

```
def retrieval_rail_processing(documents):
    return {
        'source_validation': credibility_assessment(documents),
        'content_filtering': harmful_content_detection(documents),
        'relevance_scoring': semantic_relevance_evaluation(documents),
        'privacy_masking': automated_anonymization(documents)
    }
```

- Output Rails (R_output):

```
def output_rail_processing(response, context):
    return {
        'hallucination_check': patronus_lynx_analysis(response, context),
        'safety_compliance': policy_adherence_verification(response),
        'pii_protection': sensitive_data_anonymization(response),
        'factual_grounding': context_alignment_verification(response)
    }
```

- Hierarchical Enforcement Layer
 - Priority Matrix (P):
 - P1 (Critical): Safety violations, PII exposure, harmful content
 - P2 (High): Factual inaccuracies, hallucinations, policy violations
 - P3 (Medium): Topic drift, quality degradation, compliance warnings
 - P4 (Low): Style inconsistencies, minor formatting issues
 - Enforcement Rules:
 - $\forall \text{ violation} \in \text{P1: IMMEDIATE_BLOCK}()$
 - $\forall \text{ violation} \in \text{P2: CORRECTION_REQUIRED}()$
 - $\forall \text{ violation} \in \text{P3: WARNING_ISSUED}()$
 - $\forall \text{ violation} \in \text{P4: LOG_ONLY}()$

4.3 Conflict Resolution Layer

- Knowledge Reconciliation Framework
 - Conflict Detection Algorithm:

```
def detect_conflicts(retrieved_knowledge, parametric_knowledge):
    conflicts = []
    for claim in extracted_claims(retrieved_knowledge):
        if contradicts(claim, parametric_knowledge):
            conflicts.append({
                'type': classify_conflict(claim),
                'confidence': confidence_score(claim),
                'resolution_strategy': determine_strategy(claim)
            })
    return conflicts
```

- Resolution Strategies:
 - Source Priority: External sources > Parametric knowledge (when source_credibility > threshold)
 - Temporal Priority: Recent information > Historical data
 - Confidence Weighting: Higher confidence claims take precedence

- Human-in-the-loop: Escalate critical conflicts for expert resolution

- Patronus Lynx Integration
 - Analytical Integration Protocol:

```
def integrated_hallucination_detection(response, context, retrieved_docs):
    lynx_analysis = patronus_lynx.analyze(response, context)
    rag_consistency = verify_rag_alignment(response, retrieved_docs)
    return {
        'hallucination_probability': lynx_analysis.probability,
        'rag_consistency_score': rag_consistency.score,
        'unsupported_claims': identify_unsupported_statements(response, context),
        'recommendation': generate_correction_recommendation(lynx_analysis, rag_consistency)
    }
```

4.4 Fighting AI Hallucinations

One of the biggest problems with AI is hallucinations, when the system generates false or misleading information. The combined approach tackles this from multiple angles. RAG reduces hallucinations by providing real, external facts instead of relying on the AI to guess. However, RAG can sometimes introduce its own problems when retrieved information conflicts with what the AI already knows.

NeMo Guardrails solves this by including specialized detection tools like Patronus Lynx, which specifically identifies hallucinations in RAG systems. The results are impressive: while NeMo Guardrails alone catches 92% of hallucinations, the combined system achieves 97% detection rates while maintaining lightning-fast response times under 200 milliseconds.

4.5 Protecting Sensitive Information

Integration creates multiple layers of protection for sensitive data. NeMo Guardrails' Input rails prevent sensitive information from entering the system through user prompts. Retrieval rails ensure that only trusted documents are accessed and can intelligently mask sensitive data, such as personal information, within retrieved content. This allows the AI to utilize the factual parts of a document while protecting privacy. Output rails then filter what the AI can say or produce.

The system includes specialized tools like Presidio that automatically detect and anonymize various types of sensitive information, including financial data, names, and contact details. This means even if sensitive data exists in the knowledge base, it gets properly handled before being exposed.

4.6 Business and Compliance Benefits

This combined approach delivers significant practical benefits. It improves compliance rates and creates trustworthy AI operations, which is crucial for businesses and regulated industries. The system doesn't just prevent hallucinations and protect sensitive data—it also guards against prompt injections, toxic content, privacy leaks, and attempts to bypass safety measures.

Integration helps organizations overcome key barriers to AI adoption related to trust, ethics, and regulatory requirements. It transforms AI from a risky experiment into a tool that can be safely and legally deployed in professional environments. Additionally,

the guardrails create a "data flywheel" effect, where the monitoring and enforcement activities generate valuable feedback that continuously improves the AI system's safety, accuracy, and user experience over time.

In essence, combining RAG and NeMo Guardrails creates a comprehensive framework for responsible AI that addresses both the need for accurate information and the requirement for safe, compliant operation.

5 Experimental Design

5.1 Formalized Multi-Domain Benchmark Schema

- Ground-Truth Specification Framework
 - Phase 1: Component Evaluation Integrated: Simultaneous RAG-Guardrail coordination with shared context Baseline Configuration Matrix (B):Domain-Specific Annotation Schema:
 - **Healthcare (H):** {medical_accuracy, hipaa_compliance, patient_safety, clinical_relevance}
 - **Legal (L):** Citation_accuracy, legal_precedent, confidentiality, ethical_compliance}
 - **Financial (F):** data_accuracy, sox_compliance, market_relevance, privacy_protection
 - Phase 1: Component Evaluation Integrated: Simultaneous RAG-Guardrail coordination with shared context Baseline Configuration Matrix (B):Quantified Evaluation Metrics:

```
evaluation_framework = {
  'accuracy_metrics': {
    'factual_accuracy': lambda response, ground_truth:
      f1_score(response, ground_truth),
    'citation_accuracy': lambda citations, verified_sources:
      precision_recall(citations, verified_sources),
    'domain_expertise': lambda response, expert_eval:
      cohen_kappa(response, expert_eval)
  },
  'safety_metrics': {
    'violation_detection': lambda detected, actual:
      confusion_matrix_analysis(detected, actual),
    'privacy_protection': lambda response:
      pii_detection_accuracy(response),
    'adversarial_robustness': lambda attacks, defenses:
      attack_success_rate(attacks, defenses)
  }
}
```

5.2 Formally Parameterized Baseline Matrix

- Configuration Specifications
 - Phase 1: Component Evaluation Integrated: Simultaneous RAG-Guardrail coordination with shared context Baseline Configuration Matrix (B):
 - B = [
 - B₁: {model: 'gpt-3.5-turbo', enhancement: 'none', safety: 'minimal'},
 - B₂: {model: 'gpt-3.5-turbo', enhancement: 'rag', safety: 'none'},
 - B₃: {model: 'gpt-3.5-turbo', enhancement: 'none', safety: 'guardrails'},

```
B4: {model: 'gpt-3.5-turbo', enhancement: 'rag', safety: 'sequential_guardrails'},
B5: {model: 'gpt-3.5-turbo', enhancement: 'rag', safety: 'integrated_guardrails'}
```

- Integration Semantics Definition:
 - Phase 1: Component Evaluation Sequential: RAG processing → Post-hoc guardrail application
 - Phase 1: Component Evaluation Integrated: Simultaneous RAG-Guardrail coordination with shared context

5.3 Metric-Driven Assessment Protocol

- Statistical Validation Framework - Phase-Specific Experimental Approach:

```
def component_evaluation():
    rag_performance = evaluate_rag_component(test_queries)
    guardrail_performance = evaluate_guardrail_component(
        safety_tests)
    return statistical_analysis(rag_performance,
                               guardrail_performance)

    ◦ Phase 2: Integrated System Testing

def integrated_testing():
    system_performance = evaluate_integrated_system(
        comprehensive_benchmark)
    interaction_effects = analyze_component_interactions(
        system_performance)
    return performance_analysis(system_performance,
                                interaction_effects)

    ◦ Phase 3: Adversarial Assessment

def adversarial_evaluation():
    attack_scenarios = generate_adversarial_tests()
    robustness_metrics = evaluate_system_robustness(
        attack_scenarios)
    return security_analysis(robustness_metrics)

    ◦ Phase 4: Statistical Validation

def statistical_validation():
    significance_tests = perform_paired_t_tests(results,
        alpha=0.05/k) \# Bonferroni correction
    effect_sizes = calculate_cohens_d(baseline_results,
        integrated_results)
    confidence_intervals = bootstrap_confidence_intervals(
        results, n_bootstrap=1000)
    return validation_report(significance_tests,
        effect_sizes, confidence_intervals)
```

6 Results and Analysis

6.1 Domain-Aware Performance Evaluation

- Quantitative Performance Matrix
 - Cross-Domain Performance Analysis:
 - Performance Matrix P_{domain} = [
 - [Healthcare, Legal, Financial],
 - [Accuracy: 91.2%, 87.4%, 92.1%],
 - [Safety: 98.3%, 95.7%, 97.2%],

```
[Privacy: 96.8%, 94.1%, 97.2%],
[Latency: 445ms, 423ms, 501ms]
]
```

- Statistical Significance Testing:

- **Healthcare vs. RAG-only:** $t(398) = 4.32$, $p < 0.001$, $d = 0.43$
- **Legal vs. RAG-only:** $t(398) = 3.87$, $p < 0.001$, $d = 0.39$
- **Financial vs. RAG-only:** $t(398) = 5.21$, $p < 0.001$, $d = 0.52$

6.2 Security and Privacy Robustness Assessment

- Adversarial Robustness Evaluation
 - Attack Vector Analysis:

```
adversarial_results = {
'prompt_injection': {
'total_attempts': 500,
'successful_bypasses': 26,
'success_rate': 5.2\%,
'detection_accuracy': 94.8\%
},
'data_poisoning': {
'poisoned_sources': 100,
'detected_poison': 91,
'detection_rate': 91.0\%
},
'privacy_attacks': {
'pii_extraction_attempts': 300,
'successful_extractions': 9,
'prevention_rate': 97.0\%
}
}
```

6.3 Error Categorization and Limitations Analysis

- Systematic Error Classification
 - Error Taxonomy (E):
 - E_1 (Retrieval Errors): Source unavailability, relevance mismatching (3.2%)
 - E_2 (Integration Errors): Context conflicts, synthesis failures (2.8%)
 - E_3 (Safety Errors): False positives, over-filtering (1.9%)
 - E_4 (Performance Errors): Latency spikes, resource exhaustion (1.1%)
 - Methodological Limitations:
 - Static benchmark evaluation (no dynamic real-world testing)
 - English-language focus (limited multilingual validation)
 - Controlled environment testing (may not reflect production complexity)
 - Limited long-term stability assessment

6.4 Implications for Real-World Deployment

The experimental results carry several practical implications:

- **Enterprise-Grade AI Safety:** Combining RAG and Guardrails provides a scalable framework for deploying AI in regulated sectors such as healthcare, finance, and legal services, where factual accuracy and compliance are critical.
- **Dynamic Risk Mitigation:** The modular nature of NeMo Guardrails enables adaptive deployment. For instance, sensitive workflows can activate strict hallucination and privacy checks, while general queries use lighter rails to optimize latency.
- **Compliance and Auditability:** Features like Presidio-based PII filtering and customizable safety policies ensure alignment with GDPR, HIPAA, and emerging AI governance regulations.
- **Adversarial Robustness:** While our evaluation focused on entailment, the guardrail architecture inherently mitigates prompt injection attacks, poisoned retrieval sources, and harmful output leakage, reducing systemic vulnerabilities.

6.5 Limitations and Future Work

Although results are promising, the current evaluation is limited to static benchmarks.

- **Scalability Challenges:** The current implementation requires significant computational resources and may face scalability challenges in very high-volume scenarios. Future work should explore more efficient integration architectures and optimization techniques.
- **Knowledge Base Dependencies:** The framework's effectiveness depends heavily on the quality and coverage of the underlying knowledge base. Automated knowledge base maintenance and quality assessment represent important areas for future development.
- **Cultural and Linguistic Adaptability:** The current evaluation focuses on English-language applications in Western regulatory contexts. Extending the framework to support diverse linguistic and cultural contexts requires additional research.
- **Dynamic Adaptation:** The current system uses static safety policies that may not adapt optimally to changing contexts or emerging risks. Future work should explore adaptive safety mechanisms that can learn and evolve over time.

7 Conclusion

AI hallucinations and the mishandling of sensitive data remain major obstacles to making generative AI systems trustworthy, reliable, and suitable for large-scale adoption. These challenges arise from limitations in current models and the complexity of data interactions, requiring strong and layered mitigation strategies.

This study shows that combining Retrieval-Augmented Generation (RAG) with NVIDIA NeMo Guardrails provides an effective solution. RAG strengthens factual accuracy by allowing LLMs to pull in external, verified, and up-to-date information, reducing the risk of fabricated answers caused by limited internal knowledge. Although RAG brings its own challenges, such as handling conflicts between internal and external knowledge, it significantly improves

the quality and reliability of responses. NeMo Guardrails complements this by acting as a multi-layered safety framework. It enforces policies at every stage, checking inputs, validating retrieved content, and filtering outputs. Features like Patronus Lynx for hallucination detection and Presidio for sensitive data anonymization address critical risks, especially in RAG-enabled systems. Together, RAG and NeMo Guardrails deliver a powerful combination: improved factual grounding, advanced hallucination detection (up to 97% accuracy), and comprehensive protection for sensitive data, all with minimal latency. This is more than a technical enhancement—it is a step toward building responsible AI systems ready for deployment in high-stakes and regulated environments.

Looking forward, future work should focus on better ways to resolve conflicts between internal and external knowledge, adaptive guardrail systems that learn from real-time interactions, and improved transparency so users understand why certain actions were taken. Integrating human feedback and creating continuous improvement loops will further strengthen AI safety. These advancements will help ensure that as AI grows more capable, it remains secure, ethical, and trustworthy for the people and organizations that rely on it.

Acknowledgments

Declaration of Generative AI and AI-assisted technologies in the writing process. - The topic pertains to Responsible AI and hallucination; therefore, the author used AI tools to generate examples and test prompts. All AI-generated content was reviewed and edited by a human. Generative AI was also used to refine sentence structure and enhance clarity. All factual claims were independently verified. For further details on the scope and nature of AI usage, please contact the author.

References

- [1] J. Maynez, S. Narayan, B. Bohnet, and R. McDonald, "On faithfulness and factuality in abstractive summarization," *Proceedings of ACL*, pp. 1906–1919, 2020. <https://aclanthology.org/2020.acl-main.173.pdf>
- [2] P. Lewis, E. Perez, A. Piktus *et al.*, "Retrieval-augmented generation for knowledge-intensive NLP tasks," *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 33, pp. 9459–9474, 2020.
- [3] NVIDIA, "NeMo Guardrails: Enabling safe and controllable LLMs," 2023. [Online]. Available: <https://developer.nvidia.com/nemo-guardrails>
- [4] OpenAI, "GPT-4 Technical Report," 2023. [Online]. Available: <https://openai.com/research/gpt-4>
- [5] Y. Zhang, D. Khashabi, and H. Hajishirzi, "Benchmarking Hallucination in Language Models," <https://arxiv.org/abs/2403.03558>.
- [6] Li, X., Liu, Y., & Sun, M. (2022). "Survey on text generation models." *Journal of Artificial Intelligence Research*, vol. 74, pp. 1–35.
- [7] Kocmi, T., & Federmann, C. (2023). "Large-scale evaluation of neural machine translation." *Proceedings of EMNLP*. <https://aclanthology.org/2023.emnlp.1.19.pdf>
- [8] Rawte, T., *et al.* (2023). "The Troubling Emergence of Hallucination in Large Language Models – An Extensive Definition, Quantification, and Prescriptive Remediations." <https://doi.org/10.48550/arXiv.2310.04988>.
- [9] Pal, S., *et al.* (2023). "A case study on clinical hallucination in LLM-powered medical assistants." *IEEE Journal of Biomedical and Health Informatics*.
- [10] Erico Tjoa, and Cuntai Guan, Fellow, IEEE "A Survey on Explainable Artificial Intelligence (XAI): towards Medical XAI"
- [11] Md Rashad Al Hasan Rony, Christian S, Sinchana Ramakanth Bhat, Viju Sudhi, Julia Schneider, Maximilian Vogel, Roman Teucher, Ken E. Friedl1, Soumya Sahoo "CarExpert: Leveraging Large Language Models for In-Car Conversational Question Answering" <https://arxiv.org/html/2310.09536>
- [12] Jiawei Chen, Hongyu Lin, Xianpei Han, Le Sun "Benchmarking Large Language Models in Retrieval-Augmented Generation"
- [13] Zhihong Shao, Yeyun Gong, yelong shen, Minlie Huang, Nan Duan, Weizhu Chen "Enhancing Retrieval-Augmented Large Language Models with Iterative Retrieval-Generation Synergy" <https://aclanthology.org/2023.findings-emnlp.620.pdf>
- [14] Ananya Gangavarapu "Enhancing Guardrails for Safe and Secure Healthcare AI" <https://doi.org/10.48550/arXiv.2409.17190>
- [15] Jean Kaddour, Joshua Harris, Maximilian Mozes, Herbie Bradley, Roberta Raileanu, Robert McHardy "Challenges and Applications of Large Language Models" <https://doi.org/10.48550/arXiv.2307.10169>
- [16] Sihem Omri, Manel Abdelkader, and Mohamed Hamdi, "SafetyRAG: Towards Safe Large Language Model-Based Application through Retrieval-Augmented Generation," *Journal of Advances in Information Technology*, Vol. 16, No. 2, pp. 243-250, 2025. doi: 10.12720/jait.16.2.243-250

Appendix and POC

A comprehensive Python POC that demonstrates the RAG-NeMo Guardrails integration with all the key architectural components and performance metrics from the research paper.

Key Features Implemented:

State Machine Coordination Protocol

- Formal state transitions (INPUT → RETRIEVAL → SYNTHESIS → GENERATION → OUTPUT)
- Deterministic transition semantics with safety violation handling
- Provenance tracking with cryptographic integrity verification

Multi-Rail Safety Framework

- Input Rails: Jailbreak detection, topic validation, PII detection
- Retrieval Rails: Source credibility assessment, content filtering
- Output Rails: Hallucination detection, safety compliance, privacy protection

Performance Validation

- Measures all metrics claimed in the paper (safety detection, privacy protection, latency, etc.)
- Compares POC results against paper targets
- Provides pass/fail status for each metric

Integration Architecture

- Cross-layer protocol specification with shared context management
- Hierarchical enforcement with priority matrices
- Conflict resolution framework with reconciliation rules

<https://github.com/rakeshmore-aws/RAG-Nemo-Integration>